

Dynamic Programming and Reinforcement Learning

Lecturer: *Emilie Kaufmann**emilie.kaufmann@inria.fr***Report Deadline: November 13, midnight**

Note: Your report should be *short* and based on the answers to questions Q1-Q4

Report and code should be sent by e-mail to emilie.kaufmann@inria.fr (one pdf file *yourname_TP1.pdf* + one archive *yourname_TP1.zip*), with [MVA 2016] in the title.

1 The One-Site Tree Cutting Problem

We would like to formalize the *tree cutting* problem and compute the strategy which maximizes the revenue. A tree *keeps growing over* time with a rate which may depend on the weather and it stops when it reaches a certain *maximum height*. At the same time the tree may get a *disease*, in which case it dies and loses all its value. When the company decides to *cut a tree*, it gains an amount of money which is proportional to the height of the tree. *Whenever a tree is cut* (or it is dead), a new tree has to be planted with a *fixed cost*. Knowing that the one unit of money loses value over time, find the optimal cutting strategy.

1.1 A Bit More Formal Definition of the Environment

- **State space:** the (discrete!) height of the tree (*constrained to a maximum height*)
- **Initial state:** the height of the tree is set *to one*
- **Action space** either cut or not the tree
- **Dynamics:**
 - **If no cut:** the tree grows up to a maximum height by a number of units which depend on the (random!) *weather*. It may also (*randomly!*) get a disease.
 - **If cut:** a new tree is planted with an initial height of one unit.
- **Reward:**
 - If no cut: a *fixed amount of maintenance cost*
 - If cut: the value of *each unit of wood times the height of the tree minus the cost of planting a new tree*.
- **Discount factor:** we assume a bank interest rate $r = 0.05$, and so discount factor is set of $\gamma = 1/(1+r)$.

1.2 Work to do

1. Formalize the problem more precisely (some decisions are of course arbitrary, such as the influence of the weather on the growth) and implement two functions:
 - (a) `tree_sim` which receives as input a state and an action and it returns the next state and the reward.
 - (b) `tree_MDP` which returns the dynamics and the reward function (in suitable structures).

Q1: Explicit the MDP and the parameter chosen to model the random effects.

Note: You may choose to use the representation proposed in the code `mainTP1.m` available online: the dynamics are represented by the “growth” matrix (that you can customize) and the sick probability. In that case, you can look into the function `tree_MDP.m` that is given. For `tree_simu.m` you may need to be able to simulate from finite distributions given by a probability vector; for this purpose you can use the given code `simu.m`.

2. **Policy evaluation:** define an arbitrary policy and evaluate it in the initial state using one RL method (Monte-Carlo or TD(0)) and one dynamic programming method (matrix inversion or Bellman operator).
 - **Q2:** If V_n denotes the value function computed by the RL method based on n trajectories, chart $\frac{1}{n} \sum_{k=1}^n (V_n(x_0) - V^\pi(x_0))$, where x_0 is the initial state and V^π is the value function computed with DP.

$$V_n(x_0) = \frac{1}{n} \sum_{k=1}^n \left[\sum_{t=1}^{T_{\max}} \gamma^{t-1} R_t^{(k)} \right]$$

where $(R_t^{(k)})$ is the sequence of rewards obtained when simulating the k -th trajectory (using `tree_simu`).

3. Optimal policy:
 - **Q3:** Compute the optimal policy with the two dynamic programming method seen in class, Policy Iteration and Value Iteration.
Recall that both VI and PI can be implemented using the Q-value function associated to a value function V , defined by

$$Q(x, a) = r(x, a) + \gamma \sum_{y \in \mathcal{X}} p(y|x, a) V(y).$$
 - **Q4:** For both methods, plot $\|V^* - V_k\|_\infty$ as a function of iteration k to compare the speed of convergence and discuss the relative merits of the two approaches.
For Policy Iteration, $V_k = V^{\pi_k}$, where π_k is the policy obtained after k iterations.

In the next session, we will implement Q-Learning for this problem (i.e. learning the optimal policy with a Reinforcement Learning approach)

2 Going further

1. Study how the obtained results change when changing some of the parameters of the problem (initial height, cost of planting a new tree, gain in selling a tree, and so on).
2. Consider the case where we have two sites where we can grow trees. At each point in time, the decision is whether to cut a tree and which one and the state should consider both sites. Implement the extension or discuss how it could be implemented.
3. Propose a model (and test Q-learning on it) to solve the problem sketched here <http://stackoverflow.com/questions/8337417/markov-decision-process-value-iteration-how-does-it-work>